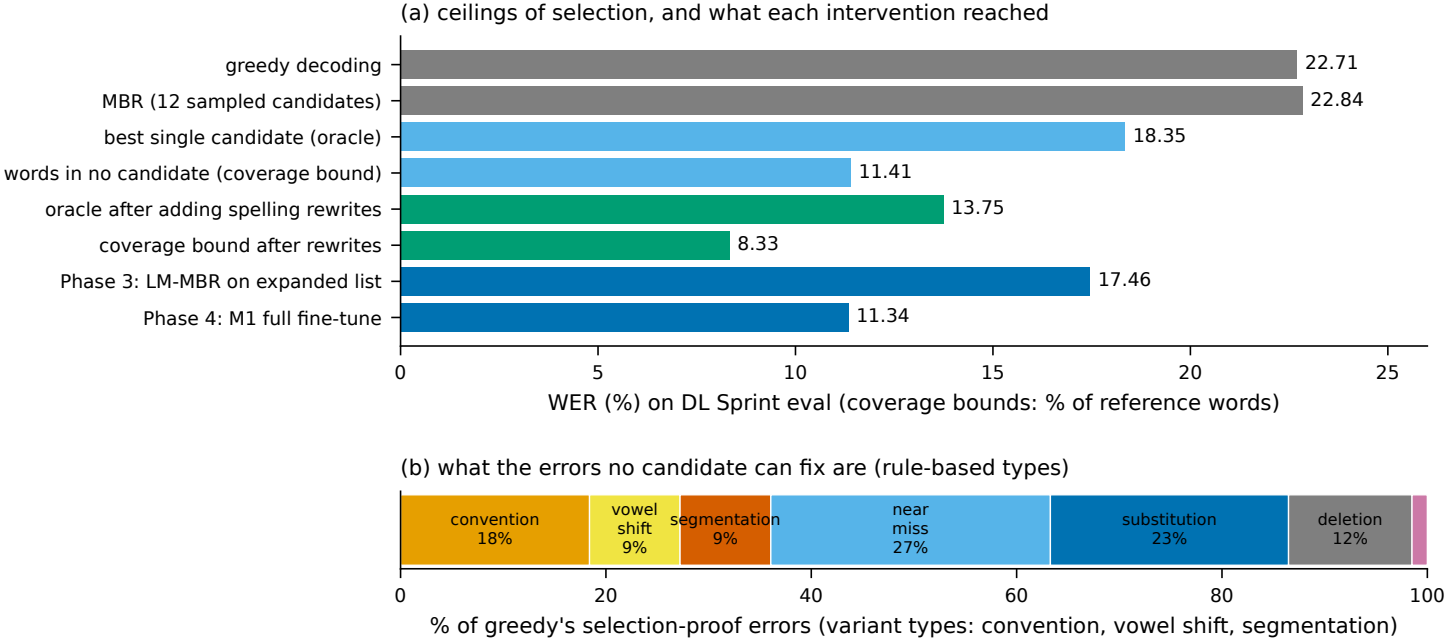
(a) ceilings of selection, and what each intervention reached
greedy decoding
MBR (12 sampled candidates)
best single candidate (oracle)
words in no candidate (coverage bound)
oracle after adding spelling rewrites
coverage bound after rewrites
Phase 3: LM-MBR on expanded list
Phase 4: M1 full fine-tune
22.71
22.84
18.35
11.41
13.75
8.33
17.46
11.34
0
5
10
15
20
25
WER (%) on DL Sprint eval (coverage bounds: % of reference words)
(b) what the errors no candidate can fix are (rule-based types)
convention 18%
vowel shift 9%
segmentation 9%
near miss 27%
substitution 23%
deletion 12%
0
20
40
60
80
100
% of greedy's selection-proof errors (variant types: convention, vowel shift, segmentation)